

CS35L Software Construction Laboratory

Lab 5: Sneha Shankar
Week 7; Lecture 1



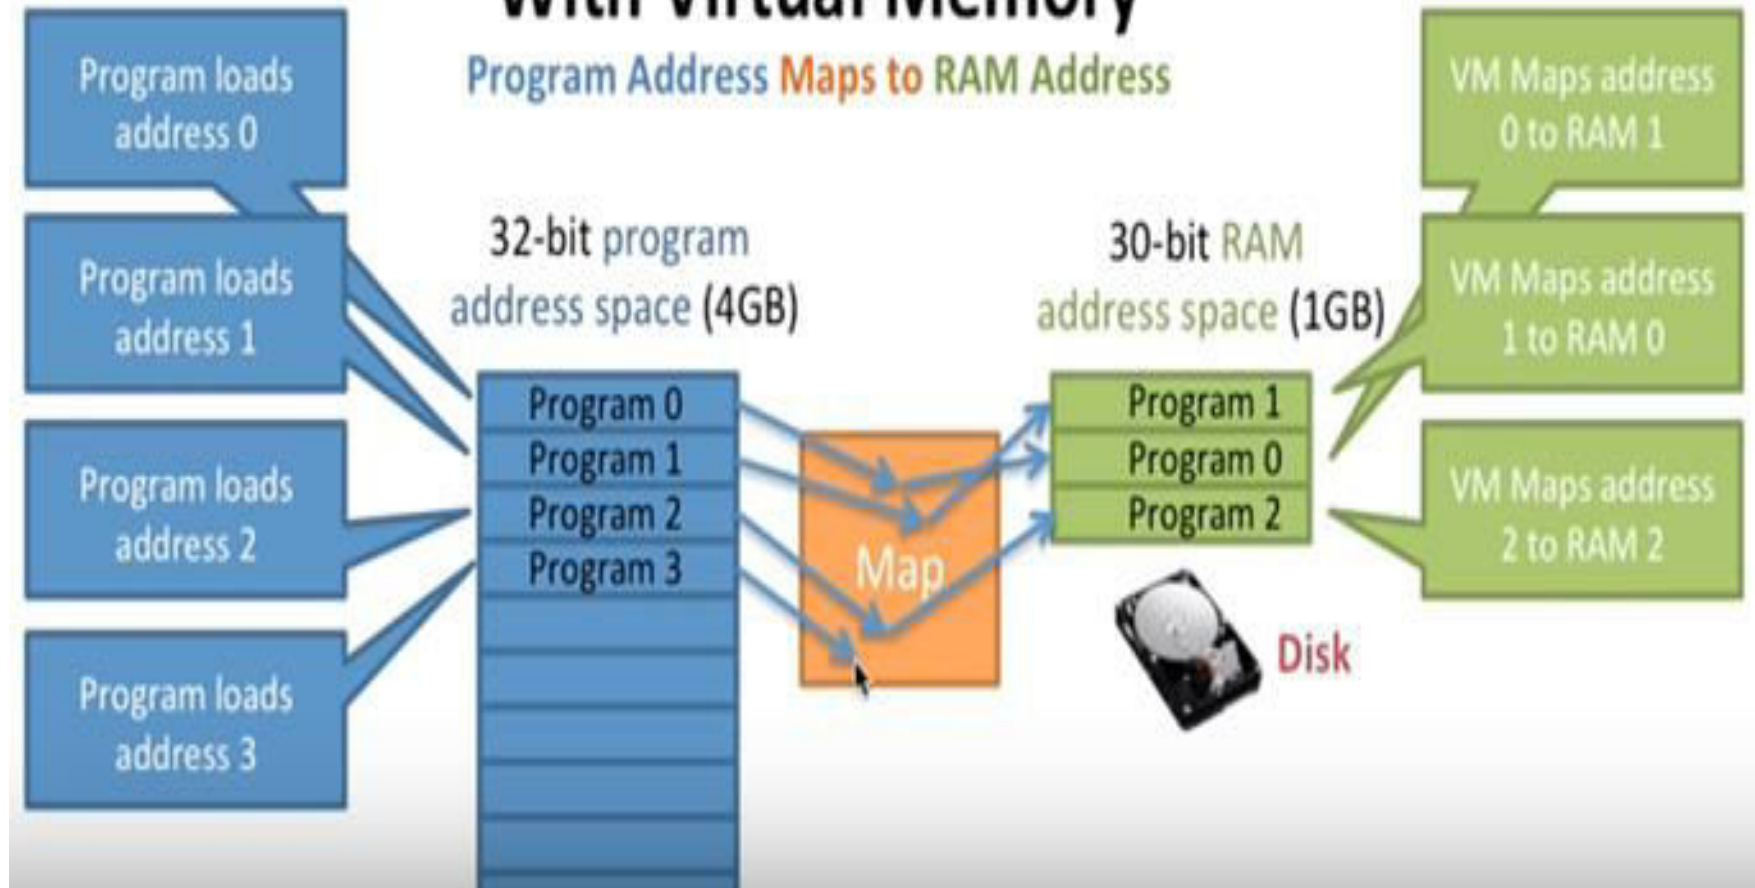
Parallelism & Multithreading

Types of resources

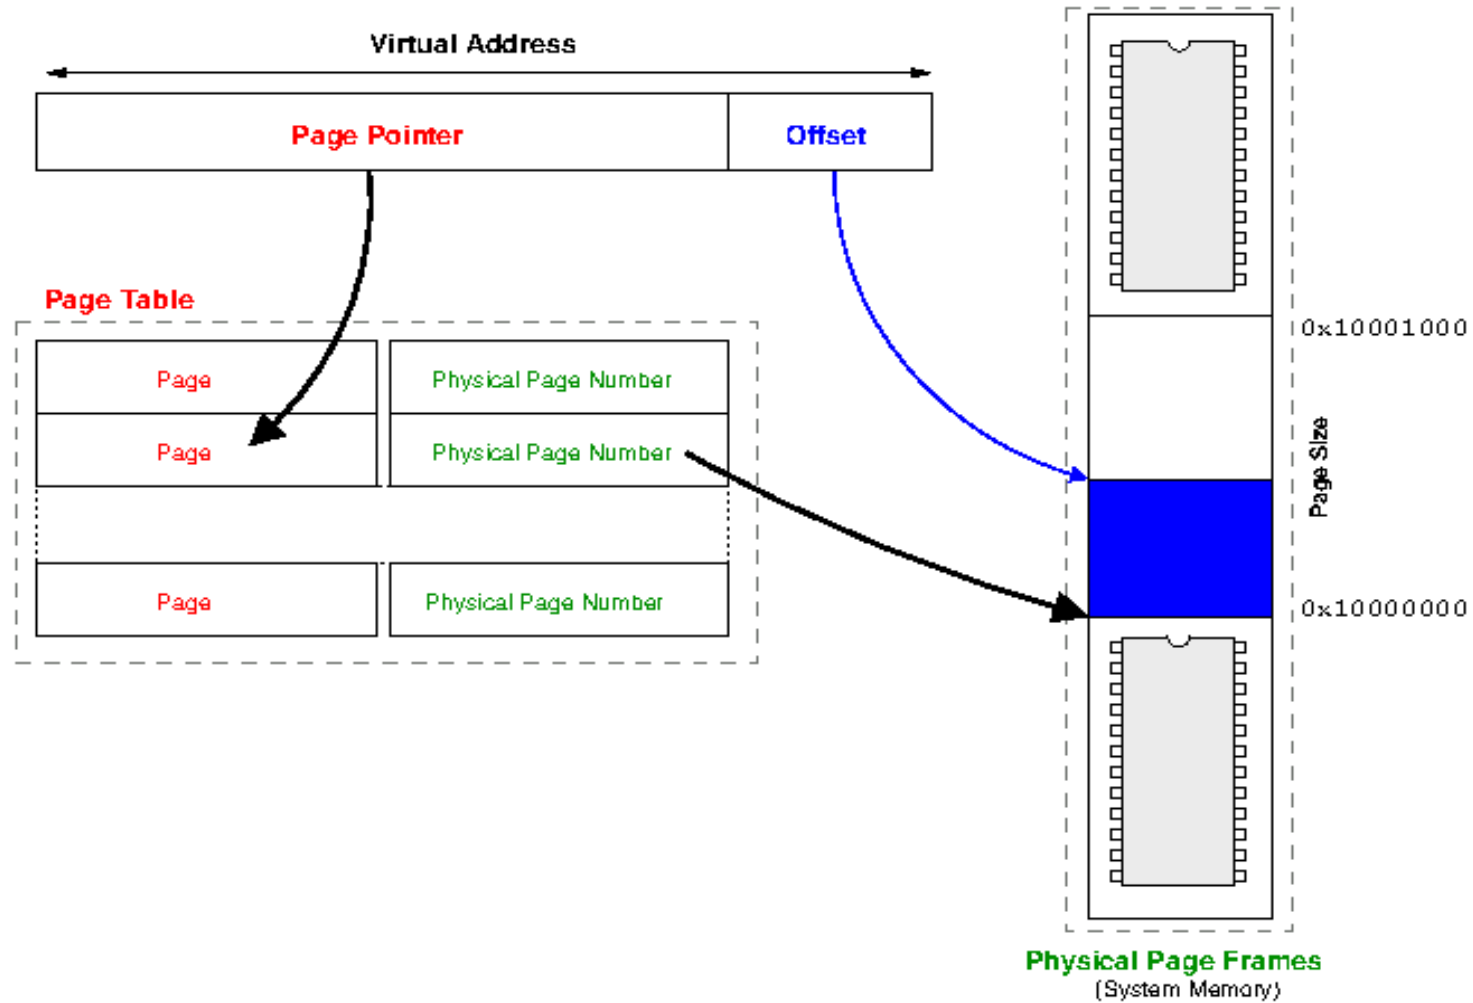
- CPU
 - It is an active resource
 - Can be used by only one runtime entity
 - Can be multiplexed in time (time sharing)
- Memory
 - Passive resource
 - Can be shared among multiple runtime entities
 - Can be multiplexed in space (allocated)

With Virtual Memory

Program Address **Maps to** RAM Address

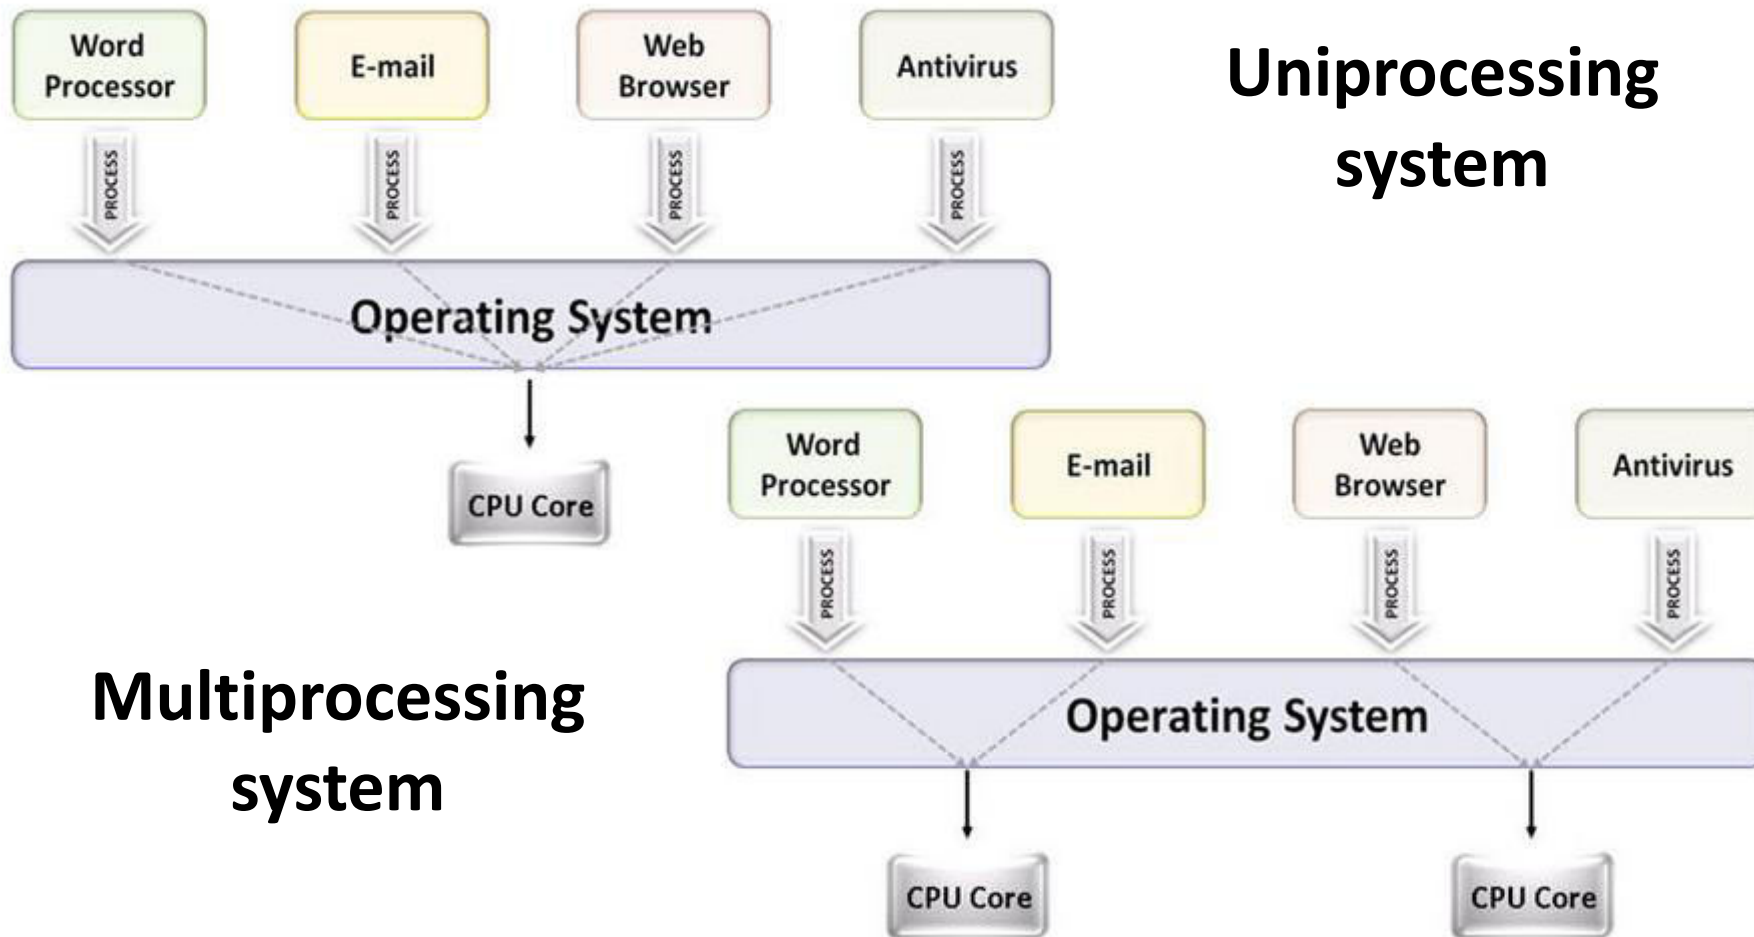


Virtual addressing



Multiprocessing

- The use of multiple CPUs/cores to run multiple tasks simultaneously



Parallelism

- Executing several computations simultaneously to gain performance
- Different forms of parallelism
 - **Multitasking**
 - Several processes are scheduled alternately or possibly simultaneously on a multiprocessing system
 - **Multithreading**
 - Same job is broken logically into pieces (threads) which may be executed simultaneously on a multiprocessing system

Parallelism

Two different parallel architectures :

- SIMD (Single Instruction Multiple Data)
 - Performs single identical action simultaneously on multiple data pieces
 - Ex. retrieving multiple files at the same time
 - Should be synchronized
- MIMD (Multiple Instruction Multiple Data)
 - Performs multiple actions simultaneously
 - Ex. addition, multiplication simultaneously
 - May/may not be synchronized

SIMD vs MIMD

- SIMD used for problems which require a lot of computation with processors performing same operation in parallel
- MIMD used for problems that break down into separate and independent parts with each part assigned to different processor simultaneously

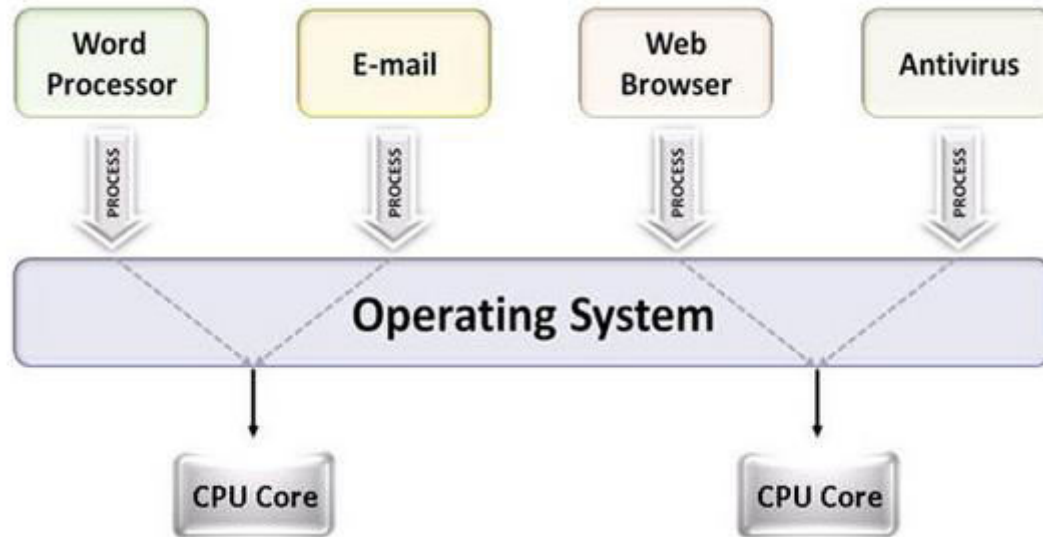
What is a thread?

- A flow of instructions, path of execution within a process
- The smallest unit of processing scheduled by OS
- A process consists of at least one thread
- Multiple threads can be run on:
 - **A uniprocessor (time-sharing)**
 - Processor switches between different threads
 - Parallelism is an illusion
 - **A multiprocessor**
 - Multiple processors or cores run the threads at the same time
 - True parallelism

Process vs Threads

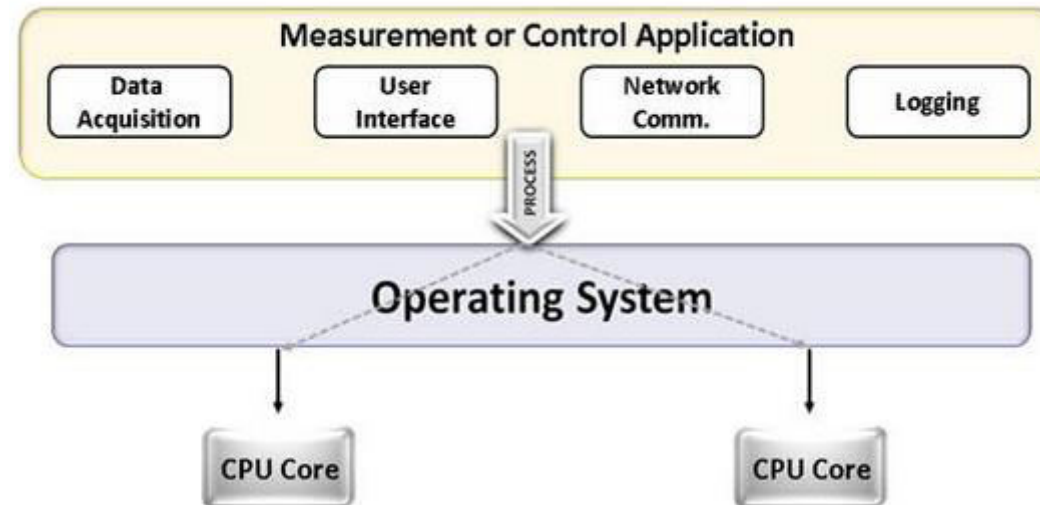
- Different processes see separate address spaces
 - good for protection, bad for sharing
- All threads in the same process share the same memory (except stack)
 - good for sharing, bad for protection
 - each thread can access the data of other thread

Multitasking vs. Multithreading



Multitasking


Multithreading



Multithreading & Multitasking: Comparison

- **Multithreading**
 - Threads share the same address space
 - Light-weight creation/destruction
 - Easy inter-thread communication
 - An error in one thread can bring down all threads in process
- **Multitasking**
 - Processes are insulated from each other
 - Expensive creation/destruction
 - Expensive IPC
 - An error in one process cannot bring down another process


Multiple
threads
sharing a
single CPU



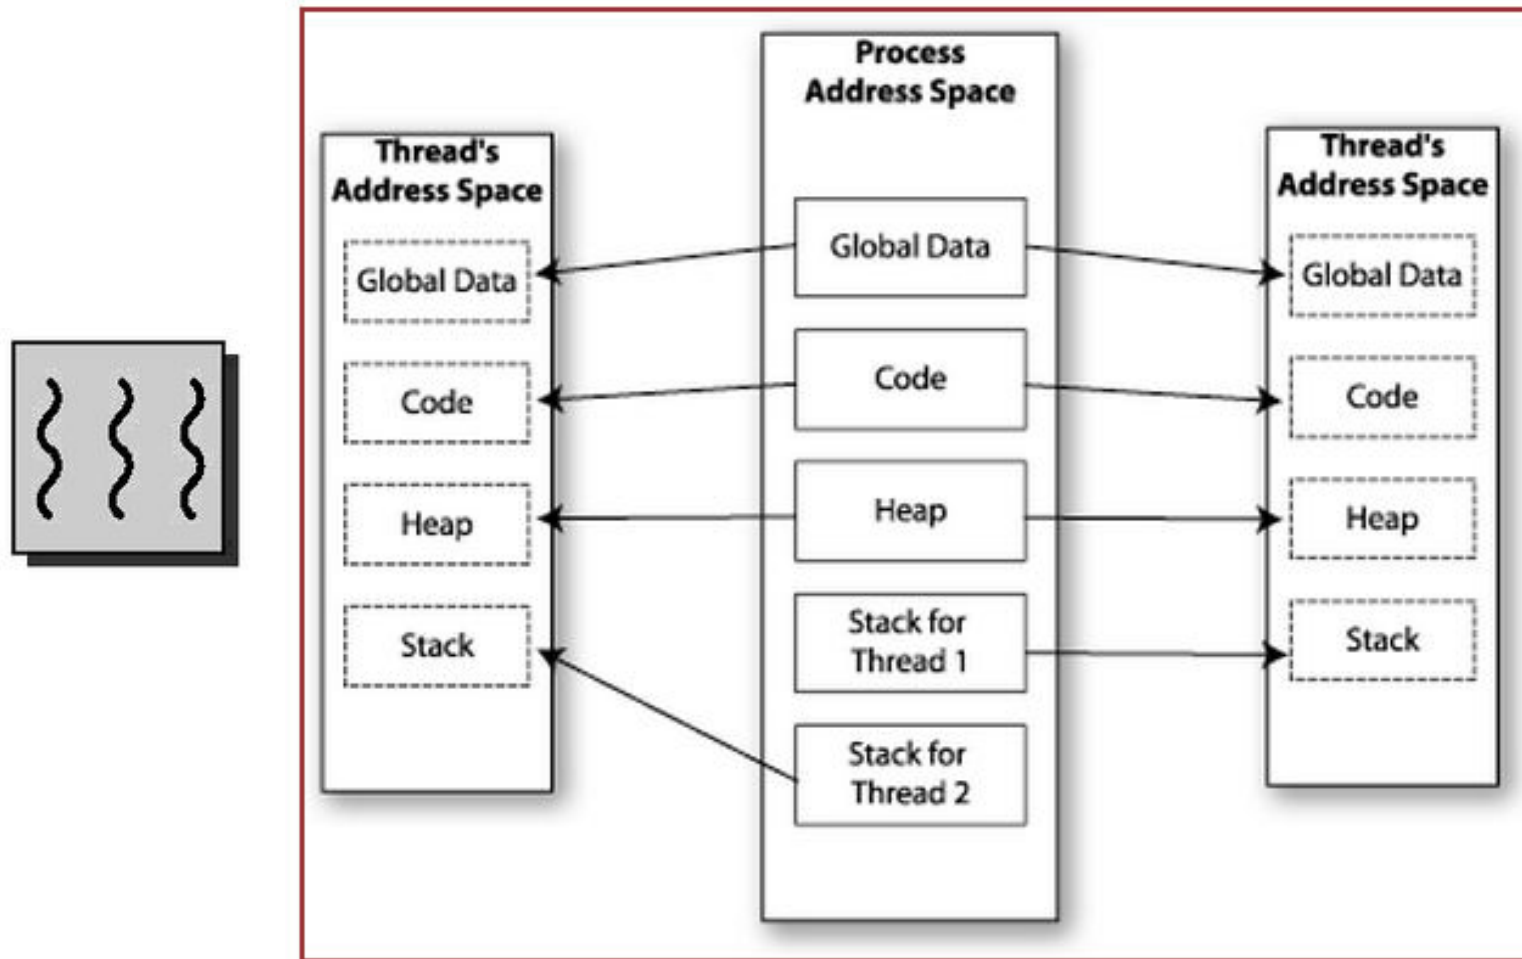
Multiple
threads on
multiple
CPUs



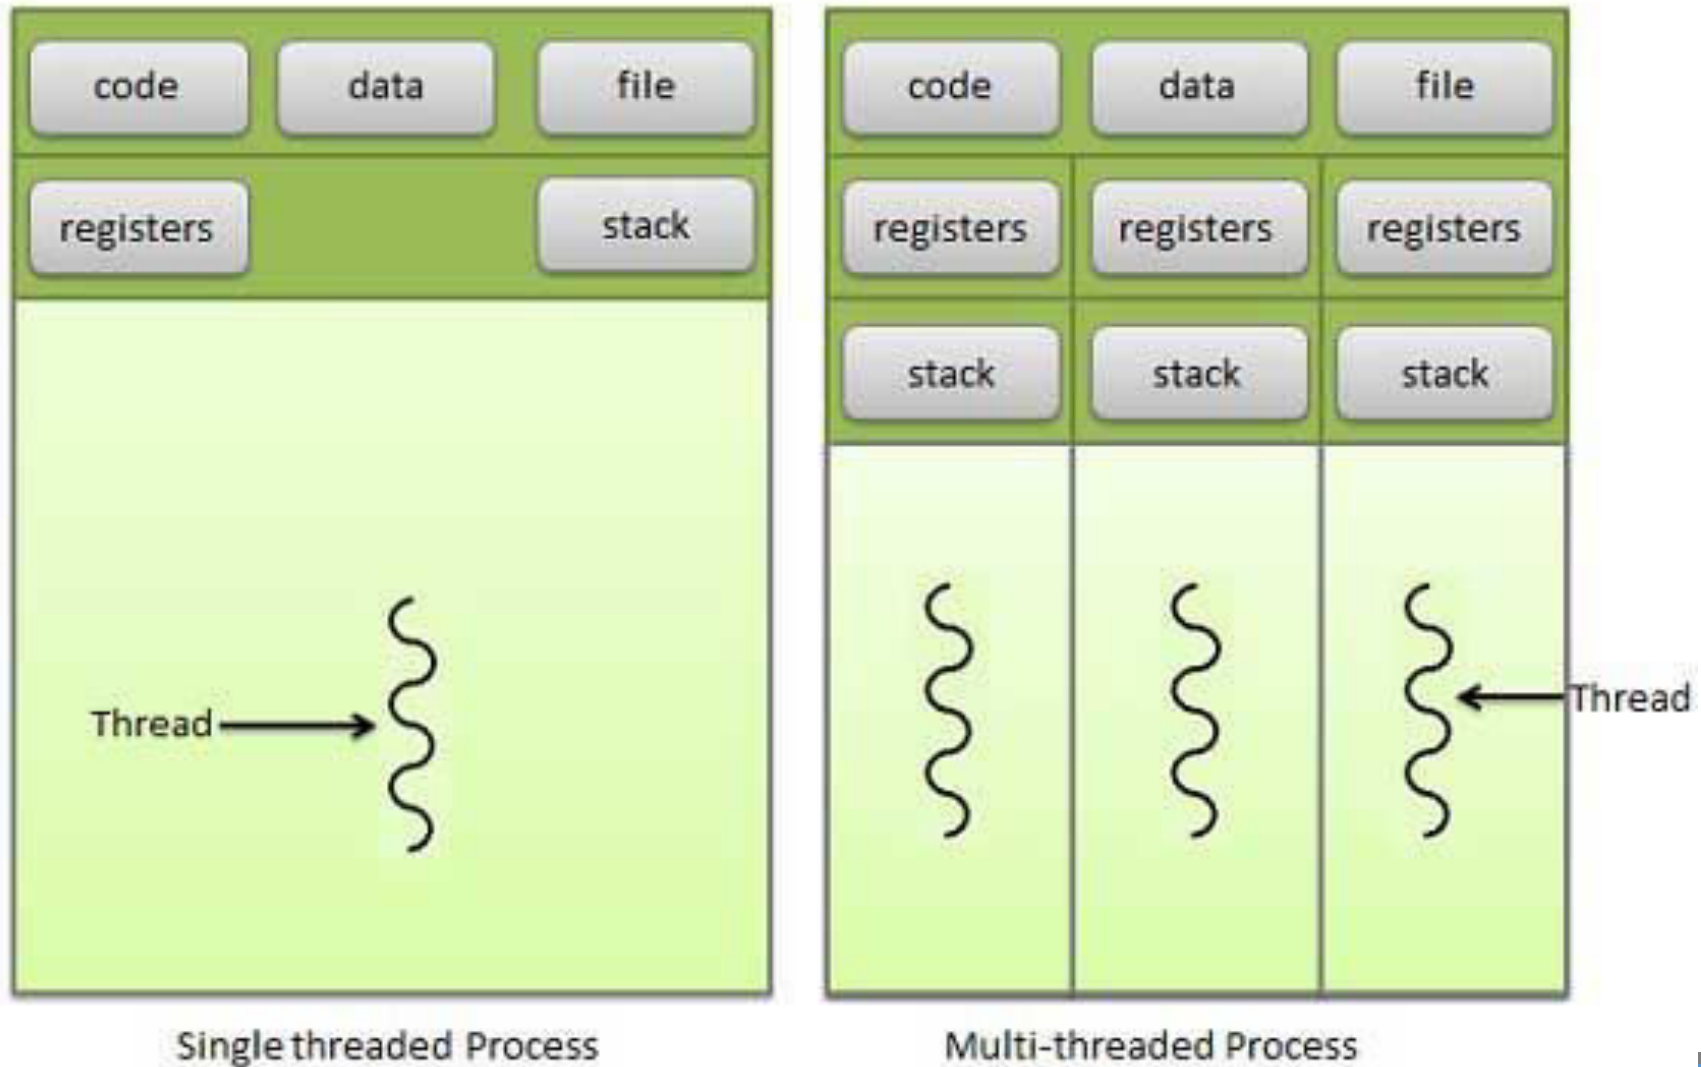
Memory Layout: Single-Threaded Program



Memory Layout: Multithreaded Program



Multithreading memory layout



Shared Memory

- Makes multithreaded programming
 - **Powerful**
 - can easily access data and share it among threads
 - **More efficient**
 - No need for system calls when sharing data
 - Thread creation and destruction less expensive than process creation and destruction
 - **Non-trivial**
 - Have to prevent several threads from accessing and changing the same shared data at the same time (synchronization)

Process/thread synchronization


Why is it needed?

- Because threads share the same resources, we need synchronization
- To prevent inconsistency

Race Condition

```
int count = 0;  
void increment()  
{  
    count = count + 1;  
}
```

Result depends on order of execution
=> Synchronization needed



Example

a=10

```
P(){  
    read(a);  
    a = a+1;  
    write(a);  
}
```

#1: P1 should first execute P and then P2.

Ans = 12

#2: P1 -> reads a=10, context switch to P2

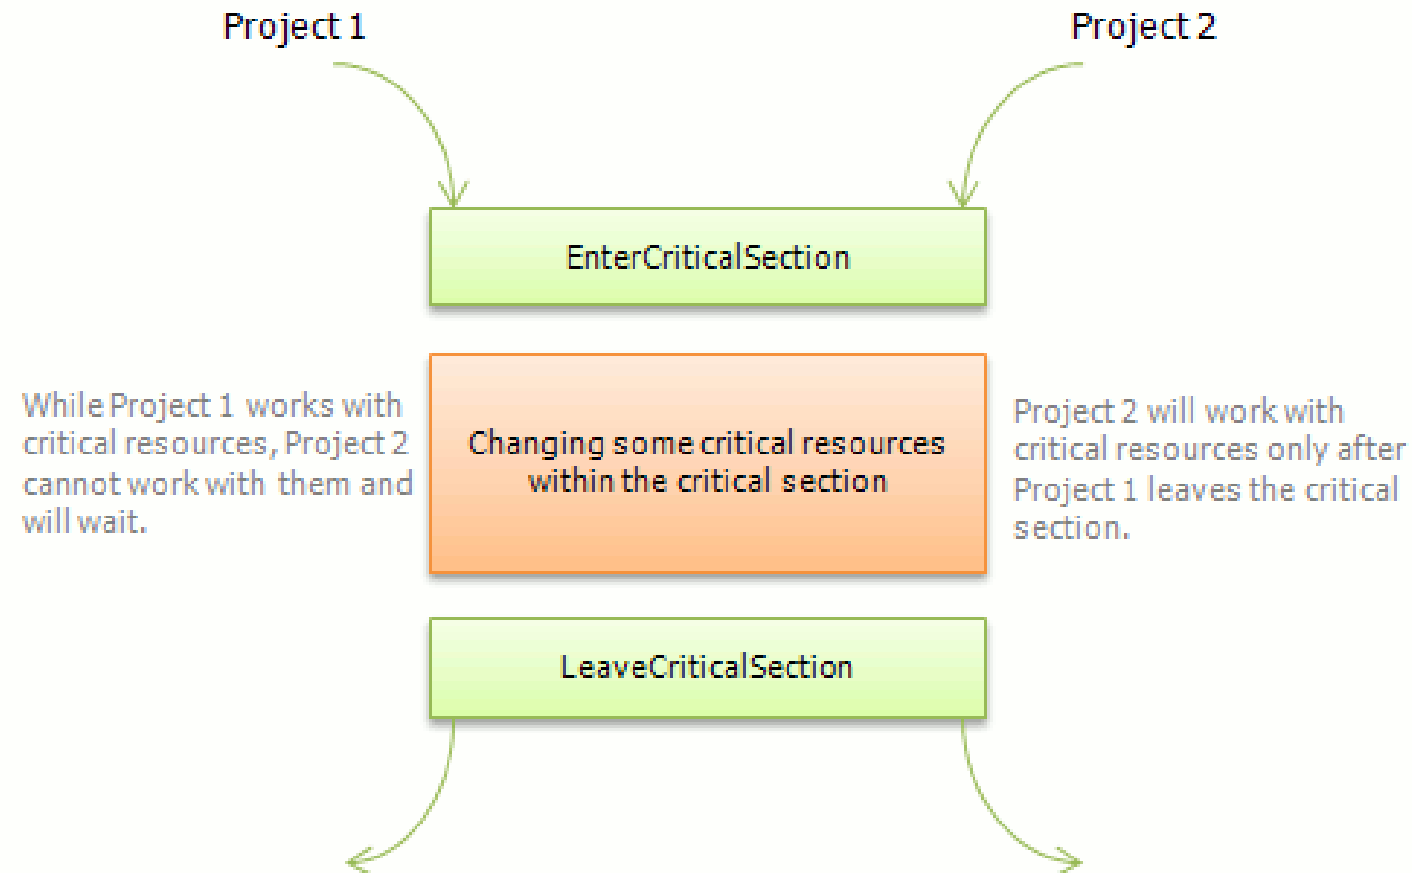
P2 -> reads a=10, adds 1 to a -> 11, switch to P1

P1 -> (has already read 10) a=11

Ans = 11

How to deal with it?

Critical section (prevents race condition)



Mutex

- Mutex is an object which allows only one thread into a critical section
- Mutex is owned by a thread
- It forces other threads which attempt to gain access to that section, to wait until the first thread has exited from the section
- Each resource has a mutex




Thread-1

Lock(mutex)
Do work
Unlock(mutex)

Thread-2

Thread-3



If thread-2 will try to access shared resource by accessing mutex object then thread-2 attempt's to take mutex will return error.



Thread-1

Thread-2

Thread-3

Lock(mutex)
Do work
Unlock(mutex)

Mutex Lock

```
int a;  
mutex-lock lock-a;  
void some_procedure(){
```

```
void another_procedure(){
```

```
    //manipulate a
```

```
}
```

Memory

```
//manipulate a
```

```
}
```

A hand holding a black pen is pointing towards a rectangular box labeled 'Memory'. Two red arrows originate from the code snippets: one from the 'manipulate a' comment in 'some_procedure()' and another from the 'manipulate a' comment in 'another_procedure()'. Both arrows point towards the 'Memory' box, illustrating that both procedures are attempting to access the same memory location simultaneously, which is why a mutex lock is needed.

Semaphores

- A semaphore is a value in a designated place in the OS (or kernel) storage that each process can check and then change.
- signaling mechanism
- restricts/allows the number of simultaneous threads of a shared resource upto a maximum number
- threads can request access to a resource (decrements the semaphore)
- threads signal that they have finished using the resource (increments the semaphore)



$S = 2$

Thread-1



$S = 2$


Thread-1





Decrease the value and check if it is less than zero. If yes wait otherwise access resource.





Semaphores v/s mutex

- Semaphore allows multiple program threads to access the finite instance of resources.
- On the other hand, Mutex allows multiple program threads to access a single shared resource but one at a time.

Multitasking

- `$ tr -cs 'A-Za-z' '[\n*]' | sort -u | comm -23 – words`
 - Process 1 (tr)
 - Process 2 (sort)
 - Process 3 (comm)
- Each process has its own address space
- How do these processes communicate?
 - Pipes/System Calls

Multithreading

- Threads share all of the process's memory except for their stacks
- => Data sharing requires no extra work (no system calls, pipes, etc.)

POSIX Threads

- import the pthread library
Example: `#include<pthread.h>`
- Use `-pthread` while compiling
- Represented by `pthread_t` (datatype)

Basic pthread Functions

There are 5 basic pthread functions:

1. **pthread_create**: creates a new thread within a process
2. **pthread_join**: waits for another thread to terminate
3. **pthread_equal**: compares thread ids to see if they refer to the same thread
4. **pthread_self**: returns the id of the calling thread
5. **pthread_exit**: terminates the currently running thread